

Synchronization problems in distributed systems

Prof. Claudio Bettini

EveryWare Lab, Dipartimento di Informatica
Università degli Studi di Milano

1

Copyright

- Le slide di questo corso sono in parte personalizzate da quelle fornite come materiale associato ai testi consigliati (Distributed Systems: Concepts and Design, 5/e, Coulouris, Dollimore, Kindberg & Blair, Addison-Wesley, e Distributed Systems: Principles and Paradigms, A. S. Tanenbaum, M. Van Steen, Prentice Hall, 2007). Parte delle slide sono invece di proprietà del docente. Tutte le slide sono soggette a diritto d'autore e quindi non possono essere ri-distribuite senza consenso. Lo stesso vale per eventuali registrazioni o videoregistrazioni delle lezioni.
- The slides for this course are partly adapted from the ones distributed by the publisher of the suggested books for this course (Distributed Systems: Concepts and Design, 5/e, Coulouris, Dollimore, Kindberg & Blair, Addison-Wesley and Distributed Systems: Principles and Paradigms, A. S. Tanenbaum, M. Van Steen, Prentice Hall, 2007). Other slides are from the teacher of this course. All the material is subject to copyright and cannot be redistributed without consent of the copyright holder. The same holds for audio and video-recordings of the classes of this course.

2

Outline

- Physical clock synchronization
- Logical clocks
- Mutual exclusion algorithms
- Election algorithms




C. Bettini - Distributed and Pervasive Systems

3

3

Clock Synchronization

When each machine has its own clock, an event that occurred after another event may nevertheless be assigned an earlier time.




C. Bettini - Distributed and Pervasive Systems

4

4

Physical Clocks

Computation of mean solar day




C. Bettini - Distributed and Pervasive Systems

5

5

UTC = TAI + leap seconds

TAI (International Atomic Time) seconds are of constant length, unlike solar seconds. Leap seconds are introduced (once in a while) to keep our time in phase with the sun.




C. Bettini - Distributed and Pervasive Systems

6

6

Slow and fast clocks

The relation between clock time and UTC when clocks tick at different rates.



C. Bettini - Distributed and Pervasive Systems

7

7

Synchronizing Physical Clocks

C. Bettini - Distributed and Pervasive Systems

8


8

Method 1: Using GNSS

- GPS is the most used Global Navigation Satellite System (GNSS)
- GNSS satellites have atomic clocks onboard
- GNSS *satellites broadcast messages* that include timing information
- GNSS *receivers listen* to multiple satellites and use trilateration to determine their own position and UTC time deviation
- The time obtained by GNSS has an accuracy in the order of micro- or even nanoseconds (but really depends on the accuracy of the GNSS)

Global Positioning System (1)

Computing a position in a two-dimensional space.



Global Positioning System (2)

Real world facts that complicate GPS

1. It takes a while before data on a satellite's position reaches the receiver.
2. The receiver's clock is generally not in sync with that of a satellite.




C. Bettini - Distributed and Pervasive Systems

11

11

Method 2: Cristian's algorithm and Network Time Protocol



Getting the current time from a time server.

C. Bettini - Distributed and Pervasive Systems


12

12

Method 3 (Internal synchronization): The Berkeley Algorithm

It is not always necessary to align the DS node clocks with external physical clocks!

(a) The time daemon asks all the other machines for their clock values.




C. Bettini - Distributed and Pervasive Systems

13

13

The Berkeley Algorithm (2)

(b) The machines answer.




C. Bettini - Distributed and Pervasive Systems

14

14

The Berkeley Algorithm (3)

(c) The time daemon tells everyone how to adjust their clock.



C. Bettini - Distributed and Pervasive Systems

15

15

Lamport's Logical Clocks

C. Bettini - Distributed and Pervasive Systems

18

18

Logical Clocks

- In many cases it is not necessary that nodes agree on their physical clocks value
- It may be sufficient to share the knowledge of a partial order of events on the DS (this is the case of the “make” example above)
- Seminal work by Lamport (1978)



The logical clocks idea


- A way for the nodes to agree on the order in which important events in the system occur
- An event is an internal event or a message being sent/received
- An integer counter is used at each node as a logical clock. It is incremented every time an interesting event occurs
- Considering a given instant in real time, the logical clocks on the different nodes may have different values



Lamport's algorithm (1)

- If a, b are events, the expression $a \rightarrow b$ denotes the relation “ a happens before b ”. It is a transitive relation.
- $C(a)$ is the logical clock value assigned by the process where a occurs. The value of a logical clock can only increase.
- Goal: if $a \rightarrow b$ holds, then $C(a) < C(b)$ should hold

Example



$a \rightarrow b, b \rightarrow c, c \rightarrow d, d \rightarrow f, e \rightarrow f, \dots$ We cannot say $e \rightarrow b$ or $b \rightarrow e$

Lamport's algorithm (1)

- If a and b occur in the same process the goal is naturally achieved
- If a is the event of sending a message m by a node and b is the event of receiving m , then $a \rightarrow b$ holds (a message must be sent before it is received). How to ensure $C(a) < C(b)$?



Lamport's algorithm


Updating counter C_i for process P_i

1. Before executing an event (internal or message sending), P_i executes $C_i \leftarrow C_i + 1$
2. When process P_i sends a message m to P_j , it sets m 's timestamp $ts(m)$ equal to C_i after having executed the previous step
3. When P_j receives message m at C_j , adjusts its own local counter as:
 $C_j \leftarrow \max\{C_j, ts(m)\} + 1$



Lamport's algorithm (2)

Three processes, each with its own clock. The clocks run at different rates.



(a)




C. Bettini - Distributed and Pervasive Systems

25

25

Lamport's algorithm (3)

Lamport's algorithm corrects the clocks.



(b)




C. Bettini - Distributed and Pervasive Systems

26

26

Logical clock adjustment in middleware



The positioning of Lamport's logical clocks in distributed systems



C. Bettini - Distributed and Pervasive Systems

27

27

Enforcing Total Order

- The algorithm can be modified by attaching a process number to the timestamp of an event (a unique number can be assigned to each process in the DS).
- P_i timestamps event e with $C_i(e).i$ where i is the process number
- For any two timestamps $T1=(C_i(a),i)$ and $T2=(C_j(b),j)$:
 - If $C_i(a) < C_j(b)$ then $T1 < T2$
 - If $C_i(a) > C_j(b)$ then $T2 < T1$
 - If $C_i(a) = C_j(b)$ and $i < j$ then $T1 < T2$
 - If $C_i(a) = C_j(b)$ and $i > j$ then $T1 > T2$



C. Bettini - Distributed and Pervasive Systems


28

28

Example (total order)

Each event in the system has a different timestamp.

Note: Total order of timestamps does not mean that we know the actual temporal relationship between any pair of events! (e.g., we still don't know $e \rightarrow b$)




C. Bettini - Distributed and Pervasive Systems

29

29

Exercise



Assuming the only events are message send and receive, what are the clock values at events a-g?




C. Bettini - Distributed and Pervasive Systems

30

30

Application: Totally-Ordered Multicasting

- Updating a replicated database and leaving it in an inconsistent state.



Implementing Totally Ordered Multicast

- Assumptions:
 - No messages are lost
 - Messages from the same sender are received in the order they were sent
- Process P_i sends timestamped message m_i to all others. The message itself is put in a local queue i
- Any incoming message at P_j is queued in queue j , according to its timestamp, and ACKed to every other process (send and receive events for messages and acks are totally ordered with Lamport)

Implementing Totally Ordered Multicast

- P_j passes a message m_i to its application if
 - m_i is at the head of queue j AND
 - m_i has been ACKed by each of the other processes
- Observation: all processes will eventually have the same copy of the local queue, therefore, all messages are passed to application in the same order everywhere



C. Bettini - Distributed and Pervasive Systems

33

33

Algorithm applied to previous example (1)

- P_1 = agent requiring update1; P_2 = agent requiring update2; P_3 = agent managing DB1; P_4 = agent managing DB2
- Lamport algorithm is used to impose a total order on logical clocks. Suppose P_1 sends its update at t_1 and P_2 at t_2 with $t_1 < t_2$
- P_1 sends $m_1=(P_1, t_1, +100)$ to all others and puts m_1 in Q_1
- P_2 sends $m_2=(P_2, t_2, +1\%)$ to all others and puts m_2 in Q_2
- P_3 receives m_1 and puts it in Q_3 ; then sends $m_3=(P_3, t_3, \text{ack}M_1)$ to all others
- P_4 receives m_2 and puts it in Q_4 ; then sends $m_4=(P_4, t_4, \text{ack}M_2)$ to all others
- P_1 and P_2 also receive m_1 and m_2 and send the corresponding acks to everyone



C. Bettini - Distributed and Pervasive Systems

34

34

Algorithm applied to previous example (2)

- P3 receives m2 and puts it in Q3; by ordering messages according to timestamps $Q3=(m1,m2)$; then sends $m5=(P3, t5, \text{ackM2})$ to all others
- P4 receives m1 and puts it in Q4; by ordering messages according to timestamps $Q4=(m1,m2)$; then sends $m6=(P4, t6, \text{ackM1})$ to all others
- When P3 receives ackM1 from ALL the other agents it knows that the order of messages in its queue is the same for everyone and can pass the updates to the application (DB). Similarly for m2. The same will do P4 that has the same order in its queue.



C. Bettini - Distributed and Pervasive Systems

35

35

Event vs Processing Time

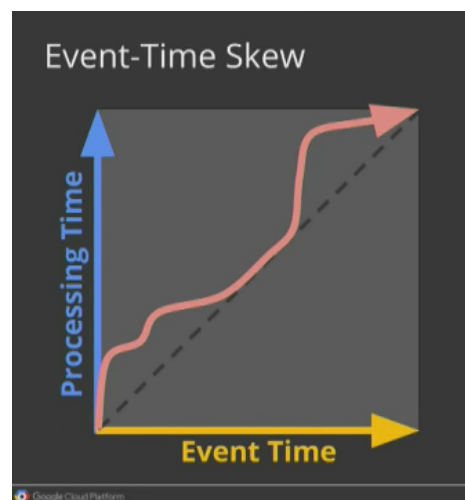
Example: Mobile gaming with teams of distributed, possibly temporarily disconnected, clients.
Goal: earning points by performing tasks. Partial score for each team?

Problem:

- events have timestamps
- game values time of answer
- unpredictable delays occur in communicating answers with their timestamps

A nice video on how Google dataflow handles the example

<https://www.youtube.com/watch?v=3UfZN59Nsk8>



C. Bettini - Distributed and Pervasive Systems

36

36

Homeworks

- Try to apply Lamport's algorithm with and without total order starting from a variation of the example in the slides.
- Consider the example of the concurrent updates of the replicated database and apply the algorithm for totally ordered multicast showing each message that is sent and the contents of the queues. Is there a problem if the transmission of an update to one of the DBs is delayed?
- Watch the video on Google stream processing tools
<https://www.youtube.com/watch?v=3UfZN59Nsk8>



C. Bettini - Distributed and Pervasive Systems

37

37



UNIVERSITÀ
DEGLI STUDI
DI MILANO

Mutual exclusion

42

Mutual exclusion

- A set of processes needs to concurrently access a shared resource
- Solution in a single node:
 - Identification of critical regions (sets of instructions that need to access a resource in exclusive mode in order to preserve consistency)
 - Use of semaphores or monitors to ensure that no other process can access the resource when the process is in its critical region.
- Solutions in distributed systems?




C. Bettini - Distributed and Pervasive Systems

53


53

A centralized algorithm


- Process 1 asks the coordinator to access the resource. Access is granted. Then, process 2 asks for access. The coordinator does not answer, but adds the request to the queue.
- When process 1 is done it notifies the coordinator that grants access to the first process in the queue in FIFO order (process 2).



(a)



(b)



(c)



C. Bettini - Distributed and Pervasive Systems

54

54

Problems of the centralized algorithm

- Single point of failure for the crash of the coordinator
- Performance bottleneck
- Difficult to distinguish a problem of the coordinator process from the unavailability of the resource



A distributed algorithm

- Ricart and Agrawala (1981)
- Assumption: total order of events and reliable message delivery (ack system)
- a) If a process P wants to use a resource R it builds a message containing:
 - 1) the name of R
 - 2) the ID of P
 - 3) the current timestampand sends the message to all the processes (including itself).



A distributed algorithm (cont.)

- b) When a process Q receives a message we have 3 cases:
- If Q is not using R and it does not require R, it answers OK to P
 - If Q is using R it does not answer and it queues the request
 - If Q wants to use R but it did not yet, it compares the message timestamp with the one in the message it sent out with its own request. The earliest wins. If the one in the message sent by P is lower it answers OK to P, otherwise it queues the message.
- c) After sending out its message process P waits for OK from all processes before accessing R.
- d) When P is done with R it sends OK to all processes in its queue and empties the queue.

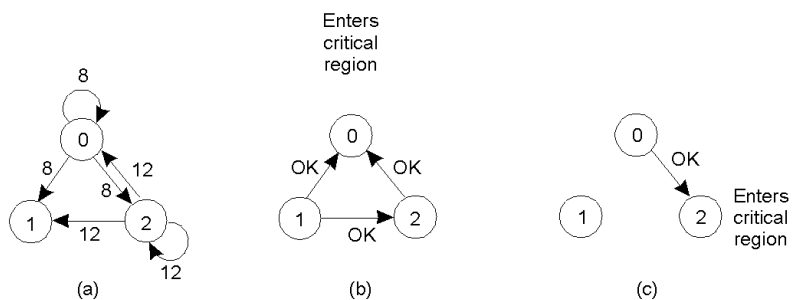


C. Bettini - Distributed and Pervasive Systems

57

57

Example



- Processes 0 and 2 want to access R at the same time; P0 broadcasts a message with timestamp 8. P2 broadcasts a message with timestamp 12
- P0 has the earliest timestamp and it wins; it adds 2 to its local queue and does not answer. P1 sends OK to both P0 and P2, since it is not interested. P2 compares timestamp 8 with its own (12) and sends OK
- When P0 is done with R, it sends OK to P2 (the only one in its queue). P2, having received OK from both P1 and P0 can now access R.



C. Bettini - Distributed and Pervasive Systems

58


58

Problems of the distributed algorithm


- No answer from a process may also be due to its crash
- Involving all processes of a distributed system may be a waste of resources

A ring algorithm (1)

- Unordered group of processes;
- A software defined logic ring.



(a)



(b)

A ring algorithm (2)

- a) At start, the token is given to process 0.
- b) The token is sent with a message from process i to process $(i+1) \text{ MOD } n$.
- c) If a process is interested in the resource, it can access when it receives the token. When done it sends the token to the next process in the ring. It cannot use twice the same token.



C. Bettini - Distributed and Pervasive Systems

61

61

Comparison of the algorithms

Algorithm	Messages to acquire/release R	Delay to acquire R (num. of messages)	Problems
Centralized	3	2	Crash of coordinator
Distributed	$2(n-1)$	$2(n-1)$	Crash of any process
Distributed as ring	1 to ∞	0 to $n-1$	token loss, crash of a process

worst case in which no process requires to acquire R



C. Bettini - Distributed and Pervasive Systems

62

62

Election algorithms

63

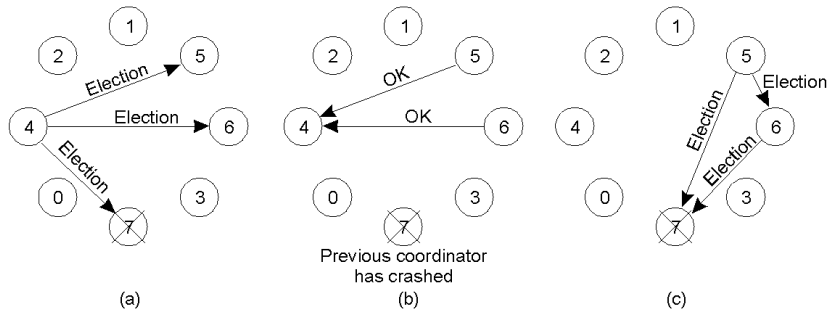
Election algorithms

- Electing a coordinator process may be useful (for example to implement a centralized solution for mutual exclusion)
- Election algorithms aim at electing the active process with the highest identifier
 - No loss of generality. For example, the process with the lowest computational load can be elected by setting $\text{Identifier}(P_i) = \langle 1/\text{load}(P_i), i \rangle$ where the unique process index i is used as a secondary ordering criteria when two processes have the same load. As another example $\text{Identifier}(P_i) = \langle \text{IP}(P_i), \text{ProcessNum}(P_i) \rangle$.
- Some algorithms assume that each process knows the IDs and can communicate with any other process. However, they cannot know if a process is currently active.

67

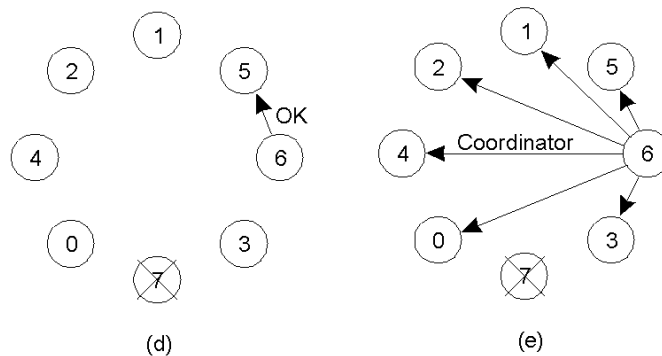
Bully Algorithm (1)

- Process 4 contacts the old coordinator; it understands that it is not active anymore and starts an election by sending a message to all processes with an ID greater than their own
- Processes 5 and 6 answer telling 4 that they will take care
- Both 5 and 6 start an election



Bully Algorithm (2)

- Process 6 tells 5 that it will take care of the election
- Process 6 wins and tells everybody



A ring-based election (1)

Chang and Roberts algorithm (1979).

Assumptions:

- N processes are logically ordered in a ring. Process P_k has a communication channel to $P_{(k+1) \bmod N}$
- Messages circulate clockwise in the ring without failures
- Processes have unique IDs.

Goal: Electing the active process with the largest ID. It must be unique (even in the case multiple elections are concurrently held)



C. Bettini - Distributed and Pervasive Systems

71

71

A ring-based election (2)

- 1) All processes are marked as *non-participant*.
- 2) When a process P_k understands that the coordinator is not answering, it starts an election by marking itself as *participant* and sending to the next node in the ring a message $\langle \text{Election}, \text{ID}(P_k) \rangle$.



C. Bettini - Distributed and Pervasive Systems

73

73

A ring-based election (3)

- 3) When a process P_m receives $\langle \text{ELECTION}, \text{ID}(P_k) \rangle$,
 - if $\text{ID}(P_k) > \text{ID}(P_m)$ P_m forwards the message in the ring and marks itself as *participant*,
 - if $\text{ID}(P_k) < \text{ID}(P_m)$ if P_m is *non-participant* it changes to *participant* and forwards $\langle \text{Election}, \text{ID}(P_m) \rangle$. If it is already *participant* it does not forward the message.
 - If $k=m$ (the received ID is its own) then it is the coordinator: it marks itself as *non-participant* and sends $\langle \text{ELECTED}, \text{ID}(P_m) \rangle$ to the next process.
- 4) When a process P_k receives $\langle \text{ELECTED}, \text{ID}(P_m) \rangle$,
 - it marks itself as *non-participant*, stores the ID of the coordinator, and unless $k=m$ it forwards the message to the next process.




C. Bettini - Distributed and Pervasive Systems

74

74

A ring-based election (4)

Note: The election was started by process 17. The highest process identifier encountered so far is 24. Participant processes are shown in a darker color



C. Bettini - Distributed and Pervasive Systems

75

75

A ring-based election (5)

Handling Failures: Failures due to crash of nodes in the ring are handled by each process by storing, not only the address of the next process, but also a few others following it in the ring. If the communication with the next process fails, the message is sent to the first among the ones following it that is active.

Concurrent elections: The use of the participant/non-participant state helps extinguishing as soon as possible unneeded messages in concurrent elections.



C. Bettini - Distributed and Pervasive Systems

76

76

A ring-based election (6)

How many messages are exchanged in the worst case (excluding failures)?

The worst case occurs when the process with highest ID is the anti-clockwise neighbor of the one that starts the election.

In this case we need $N-1$ messages to reach the node, other N messages to conclude the election and N messages to announce the coordinator. Hence $3N-1$.



C. Bettini - Distributed and Pervasive Systems

77

77

Homework

- Simulate the execution of the Ricart and Agrawala algorithm for mutual exclusion in a system including 4 processes with 3 of them making a concurrent request for the same resource. Show for each step the messages being exchanged and the status of the queues.
- Simulate the execution of the Bully election algorithm in a system with processes P2, P6, P7, P10, P15, and P20, when process P20 (the former coordinator) crashes and P7 starts an election.
- Simulate the execution of the Chang and Roberts ring-based election algorithm for the same set of processes as in the previous exercise, considering that P2 starts an election concurrently with P7.

